

01/30/18 10:11:36 /Users/caiglencross/Documents/MachineLearning/ps2/source/titanic.py

```

1  """
2  Author      : Cai Glencross and Katie Li
3  Class      : HMC CS 158
4  Date       : 2018 Jan 30
5  Description : Titanic
6  """
7
8  # Use only the provided packages!
9  import math
10 import csv
11 from util import *
12
13 from sklearn.tree import DecisionTreeClassifier
14 from sklearn.model_selection import train_test_split
15 from sklearn import metrics
16
17 #####
18 # classes
19 #####
20
21 class Classifier(object) :
22     """
23     Classifier interface.
24     """
25
26     def fit(self, X, y):
27         raise NotImplementedError()
28
29     def predict(self, X):
30         raise NotImplementedError()
31
32
33 class MajorityVoteClassifier(Classifier) :
34
35     def __init__(self) :
36         """
37         A classifier that always predicts the majority class.
38
39         Attributes
40         -----
41         prediction_ -- majority class
42         """
43         self.prediction_ = None
44
45     def fit(self, X, y) :
46         """
47         Build a majority vote classifier from the training set (X, y).
48
49         Parameters
50         -----
51         X      -- numpy array of shape (n,d), samples
52         y      -- numpy array of shape (n,), target classes
53
54         Returns
55         -----
56         self -- an instance of self
57         """

```

```

58     vals, counts = np.unique(y, return_counts=True)
59     majority_val, majority_count = max(zip(vals, counts), key=lambda (val,
count): count)
60     self.prediction_ = majority_val
61     return self
62
63     def predict(self, X) :
64         """
65         Predict class values.
66
67         Parameters
68         -----
69         X      -- numpy array of shape (n,d), samples
70
71         Returns
72         -----
73         y      -- numpy array of shape (n,), predicted classes
74         """
75         if self.prediction_ is None :
76             raise Exception("Classifier not initialized. Perform a fit first.")
77
78         n,d = X.shape
79         y = [self.prediction_] * n
80         return y
81
82
83     class RandomClassifier(Classifier) :
84
85         def __init__(self) :
86             """
87             A classifier that predicts according to the distribution of the classes.
88
89             Attributes
90             -----
91             probabilities_ -- class distribution dict (key = class, val =
probability of class)
92             """
93             self.probabilities_ = None
94
95         def fit(self, X, y) :
96             """
97             Build a random classifier from the training set (X, y).
98
99             Parameters
100            -----
101            X      -- numpy array of shape (n,d), samples
102            y      -- numpy array of shape (n,), target classes
103
104            Returns
105            -----
106            self -- an instance of self
107            """
108
109            ### ===== TODO : START ===== ###
110
111            # part c: set self.probabilities_ according to the training set
112
113            vals, counts = np.unique(y, return_counts=True)
114            majority_val, majority_count = max(zip(vals, counts), key=lambda (val,
count): count)

```

```

115     minority_val, minority_count = min(zip(vals, counts), key=lambda (val,
count): count)
116
117     # find the highest probability
118     self.proBABILITIES_ = {majority_val: float(majority_count)/sum(counts),
119                           minority_val: float(minority_count)/sum(counts)}
120
121     ### ===== TODO : END ===== ###
122
123     return self
124
125     def predict(self, X, seed=1234) :
126         """
127         Predict class values.
128
129         Parameters
130         -----
131             X      -- numpy array of shape (n,d), samples
132             seed -- integer, random seed
133
134         Returns
135         -----
136             y      -- numpy array of shape (n,), predicted classes
137         """
138         if self.proBABILITIES_ is None :
139             raise Exception("Classifier not initialized. Perform a fit first.")
140         np.random.seed(seed)
141
142         ### ===== TODO : START ===== ###
143         # part c: predict the class for each test example
144         # hint: use np.random.choice (be careful of the parameters)
145
146         class_distr = self.proBABILITIES_
147         n, d = X.shape
148
149         # Create a random distribution based on given probabilities of survival
150         random_ys = np.random.choice(class_distr.keys(), n, p =
class_distr.values())
151         y = random_ys
152         ### ===== TODO : END ===== ###
153
154         return y
155
156
157     #####
158     # functions
159     #####
160
161     def error(clf, X, y, ntrials=100, test_size=0.2) :
162         """
163         Computes the classifier error over a random split of the data,
164         averaged over ntrials runs.
165
166         Parameters
167         -----
168             clf      -- classifier
169             X        -- numpy array of shape (n,d), features values
170             y        -- numpy array of shape (n,), target classes
171             ntrials  -- integer, number of trials
172             test_size -- float (between 0.0 and 1.0) or int,

```

```

173         if float, the proportion of the dataset to include in the
test split
174         if int, the absolute number of test samples
175
176     Returns
177     -----
178         train_error -- float, training error
179         test_error  -- float, test error
180     """
181
182     ### ===== TODO : START ===== ###
183     # part b: compute cross-validation error over ntrials
184     # hint: use train_test_split (be careful of the parameters)
185     SEED = 1234
186     np.random.seed(SEED)
187
188     train_error = 0
189     test_error = 0
190
191     test_error_total = 0
192     train_error_total = 0
193
194     n, d = X.shape
195
196
197
198     for t in range(0, ntrials):
199         # get the split of data
200         X_train = np.empty((0,d)); X_test = np.empty((0,d))
201         y_train = []; y_test = []
202
203         for i in range(0, n):
204             if (np.random.random_sample() < (1.0 - test_size)):
205                 X_train = np.vstack((X_train, X[i,:]))
206                 y_train.append(y[i])
207             else:
208                 X_test = np.vstack((X_test, X[i,:]))
209                 y_test.append(y[i])
210
211         # print "X_train shape is", X_train.shape
212         # print "shape of first element of X_train is:", X[0,:].shape
213         # print "shape of second elt:", X[1,:].shape
214         # print "shape of stack is:", np.vstack((X[0,:],X[1,:])).shape
215         # print "y train length is:", len(y)
216         clf.fit(X_train,y_train)
217         y_pred_test = clf.predict(X_test)
218         y_pred_train = clf.predict(X_train)
219         test_error_total += 1 - metrics.accuracy_score(y_pred_test, y_test,
normalize=True)
220         train_error_total += 1 - metrics.accuracy_score(y_pred_train, y_train,
normalize=True)
221
222     train_error = train_error_total/ntrials
223     test_error = test_error_total/ntrials
224     ### ===== TODO : END ===== ###
225
226     return train_error, test_error
227
228
229 def write_predictions(y_pred, filename, yname=None) :
```

```

230     """Write out predictions to csv file."""
231     out = open(filename, 'wb')
232     f = csv.writer(out)
233     if yname :
234         f.writerow([yname])
235     f.writerows(zip(y_pred))
236     out.close()
237
238
239 #####
240 # main
241 #####
242
243 def main():
244     # load Titanic dataset
245     titanic = load_data("titanic_train.csv", header=1, predict_col=0)
246     X = titanic.X; Xnames = titanic.Xnames
247     y = titanic.y; yname = titanic.yname
248     n,d = X.shape # n = number of examples, d = number of features
249
250
251
252     #=====
253     # train Majority Vote classifier on data
254     print 'Classifying using Majority Vote...'
255     clf = MajorityVoteClassifier() # create MajorityVote classifier, which
includes all model parameters
256     clf.fit(X, y) # fit training data using the classifier
257     y_pred = clf.predict(X) # take the classifier and run it on the
training data
258     train_error = 1 - metrics.accuracy_score(y, y_pred, normalize=True)
259     print '\t-- training error: %.3f' % train_error
260
261
262
263     ### ===== TODO : START ===== ###
264     # part a: evaluate training error of Decision Tree classifier
265     print 'Classifying using Decision Tree...'
266     clf = DecisionTreeClassifier(criterion="entropy")
267     clf.fit(X, y)
268     y_pred = clf.predict(X)
269     train_error = 1 - metrics.accuracy_score(y, y_pred, normalize=True)
270     print '\t-- training error: %.3f' % train_error
271
272     ### ===== TODO : END ===== ###
273
274
275
276     # note: uncomment out the following lines to output the Decision Tree graph
277     """
278     # save the classifier -- requires GraphViz and pydot
279     import StringIO, pydot
280     from sklearn import tree
281     dot_data = StringIO.StringIO()
282     tree.export_graphviz(clf, out_file=dot_data,
283                         feature_names=Xnames,
284                         class_names=["Died", "Survived"])
285     graph = pydot.graph_from_dot_data(dot_data.getvalue())
286     graph.write_pdf("dtree.pdf")
287     """

```

```

288
289
290
291     ### ===== TODO : START ===== ###
292     # part b: use cross-validation to compute average training and test error of
classifiers
293     print 'Investigating various classifiers...'
294
295     print 'MajorityVoteClassifier Error'
296     train_error, test_error = error(MajorityVoteClassifier() , X, y,
ntrials=100, test_size=0.2)
297     print '\t-- training error: %.3f' % train_error
298     print '\t-- test error: %.3f' % test_error
299
300     print 'RandomClassifier Error'
301     train_error, test_error = error(RandomClassifier() , X, y, ntrials=100,
test_size=0.2)
302     print '\t-- training error: %.3f' % train_error
303     print '\t-- test error: %.3f' % test_error
304
305     print 'DecisionTreeClassifier Error'
306     train_error, test_error = error(DecisionTreeClassifier(criterion="entropy"),
X, y, ntrials=100, test_size=0.2)
307     print '\t-- training error: %.3f' % train_error
308     print '\t-- test error: %.3f' % test_error
309
310
311
312     ### ===== TODO : END ===== ###
313
314
315
316     ### ===== TODO : START ===== ###
317     # part c: investigate decision tree classifier with various depths
318     print 'Investigating depths...'
319     clfs = [MajorityVoteClassifier(), RandomClassifier()]
320     label_clfs = ["MVC", "RC"]
321     for clf_index in range(0, len(clfs)):
322         train_error_arr = []
323         test_error_arr = []
324         depths = range(1, 21)
325         for depth in depths:
326             clf = clfs[clf_index]
327             train_error, test_error = error(clf, X, y)
328             train_error_arr.append(train_error)
329             test_error_arr.append(test_error)
330
331         plt.plot(depths, train_error_arr, label = label_clfs[clf_index] + "
train error")
332         plt.plot(depths, test_error_arr, label = label_clfs[clf_index] + " test
error")
333
334     train_error_arr = []
335     test_error_arr = []
336     # only do the DecisionTreeClassifier One
337     for depth in depths:
338         clf = DecisionTreeClassifier(criterion="entropy", max_depth=depth)
339         train_error, test_error = error(clf, X, y)
340         train_error_arr.append(train_error)
341         test_error_arr.append(test_error)

```

```

342
343 plt.plot(depths, train_error_arr, label = "DT train error")
344 plt.plot(depths, test_error_arr, label = "DT test error")
345 plt.legend()
346 plt.xlabel("Max Depth")
347 plt.ylabel("Average Error")
348 plt.show()
349 ### ===== TODO : END ===== ###
350
351
352
353
354 ### ===== TODO : START ===== ###
355 # part d: investigate decision tree classifier with various training set
sizes
356 print 'Investigating training set sizes...'
357
358 clfs = [MajorityVoteClassifier(), RandomClassifier(),
DecisionTreeClassifier(criterion="entropy", max_depth=6)]
359 label_clfs = ["MVC", "RC", "DT"]
360 split_sizes = map(lambda u: u*.05, range(1,20))
361 training_splits = map(lambda u: 1 - u, split_sizes)
362
363 for clf_index in range(0, len(clfs)):
364     train_error_arr = []
365     test_error_arr = []
366     for split in split_sizes:
367         clf = clfs[clf_index]
368         train_error, test_error = error(clf,X,y,test_size=split)
369         train_error_arr.append(train_error)
370         test_error_arr.append(test_error)
371
372     plt.plot(training_splits, train_error_arr, label = label_clfs[clf_index]
+ " train error")
373     plt.plot(training_splits, test_error_arr, label = label_clfs[clf_index]
+" test error")
374
375     plt.xlabel("Training Data Size")
376     plt.ylabel("Average Error")
377     plt.legend()
378     plt.show()
379
380
381 ### ===== TODO : END ===== ###
382
383
384
385 ### ===== TODO : START ===== ###
386 # Contest
387 # uncomment write_predictions and change the filename
388
389 # evaluate on test data
390 titanic_test = load_data("titanic_test.csv", header=1, predict_col=None)
391 X_test = titanic_test.X
392 y_pred = clf.predict(X_test) # take the trained classifier and run it on
the test data
393 #write_predictions(y_pred, "../data/yjw_titanic.csv", titanic.yname)
394
395 ### ===== TODO : END ===== ###
396 print 'Done'

```

```
397  
398 if __name__ == "__main__":  
399     main()
```