

02/12/18 10:08:56

/Users/caiglencross/Documents/MachineLearning/ps2/ps4/source/phoneme.py

```

1  """
2  Author      : Cai Glencross & Katie Li
3  Class      : HMC CS 158
4  Date       : 2018 Feb 5
5  Description : Perceptron vs Logistic Regression on a Phoneme Dataset
6  """
7
8  # utilities
9  from util import *
10
11 # scipy libraries
12 from scipy import stats
13
14 # scikit-learn libraries
15 from sklearn import preprocessing
16 from sklearn import metrics
17 from sklearn import model_selection
18 from sklearn.dummy import DummyClassifier
19 from sklearn.linear_model import Perceptron, LogisticRegression
20
21 #####
22 ##
23 # functions
24 #####
25 ##
26
27 def cv_performance(clf, train_data, kfs) :
28     """
29     Determine classifier performance across multiple trials using
30     cross-validation
31
32     Parameters
33     -----
34     clf          -- classifier
35     train_data   -- Data, training data
36     kfs          -- array of size n_trials
37                   -- each element is one fold from
38                   model_selection.KFold
39
40     Returns
41     -----
42     scores       -- numpy array of shape (n_trials, n_fold)
43                   -- each element is the (accuracy) score of one
44                   fold in one trial
45     """
46
47

```

```

42     n_trials = len(kfs)
43     n_folds = kfs[0].n_splits
44     scores = np.zeros((n_trials, n_folds))
45
46     ### ===== TODO : START ===== ###
47     # part b: run multiple trials of CV
48     for n in range(n_trials):
49         scores[n,:] = cv_performance_one_trial(clf, train_data,
kfs[n])
50
51     ### ===== TODO : END ===== ###
52
53     return scores
54
55
56 def cv_performance_one_trial(clf, train_data, kf) :
57     """
58     Compute classifier performance across multiple folds using
cross-validation
59
60     Parameters
61     -----
62         clf          -- classifier
63         train_data   -- Data, training data
64         kf           -- model_selection.KFold
65
66     Returns
67     -----
68         scores       -- numpy array of shape (n_fold, )
69                        each element is the (accuracy) score of one
fold
70     """
71
72     scores = np.zeros(kf.n_splits)
73
74     ### ===== TODO : START ===== ###
75     i = 0;
76     for train_index, test_index in kf.split(train_data.X):
77         temp_train_data = train_data.X[train_index]
78         temp_test_data = train_data.X[test_index]
79         y_train = train_data.y[train_index]
80         y_test = train_data.y[test_index]
81
82         clf.fit(temp_train_data, y_train)
83         y_pred_test = clf.predict(temp_test_data)
84         test_accuracy = metrics.accuracy_score(y_pred_test, y_test,
normalize = True)
85
86         scores[i] = test_accuracy
87         i += 1

```

```

88     ### ===== TODO : END ===== ###
89
90     return scores
91
92
93     #####
94     ##
95     #####
96     ##
97     def main() :
98         np.random.seed(1234)
99
100        #=====
101        # load data
102        train_data = load_data("phoneme_train.csv")
103
104        ### ===== TODO : START ===== ###
105        # part a: is data linearly separable?
106        clf = Perceptron()
107        clf.fit(train_data.X, train_data.y)
108        print "coefs = %s, iteration = %s" % (clf.coef_, clf.n_iter_)
109        ### ===== TODO : END ===== ###
110
111        ### ===== TODO : START ===== ###
112        # part c-d: compare classifiers
113        # make sure to use same folds across all runs
114        N_SPLITS = 10
115        N_TRIALS = 10
116
117        # Order: DummyClassifier, Perceptron, LogisticRegression
118        # generate the kfs
119        kfs = []
120        for n in range(N_TRIALS):
121            kf = model_selection.KFold(n_splits=N_SPLITS, shuffle=True)
122            kfs.append(kf)
123
124        descriptive_stats = {'mean':[], 'stdev':[]}
125        dummyclassifier = DummyClassifier()
126        perceptron = Perceptron()
127        logisticregression = LogisticRegression()
128        clfs = [dummyclassifier, perceptron, logisticregression]
129
130        score_array = []
131        for clf in clfs:
132            scores = cv_performance(clf, train_data, kfs)
133            mean_result = np.mean(scores)
134            stdev_result = np.std(scores)
135            descriptive_stats['mean'].append(mean_result)

```

```
136         descriptive_stats['stdev'].append(stdev_result)
137         score_array.append(scores)
138
139     # print out descriptive_stats
140     print "Descriptive stats in the order: Dummy, Perceptron,
Logistic"
141     print descriptive_stats
142
143     # Do the t-test
144     compare_t_tests = [[0,1],[0,2],[1,2]]
145
146     pvalues = []
147     for combo in compare_t_tests:
148         result = stats.ttest_rel(score_array[combo[0]].flatten(),
score_array[combo[1]].flatten())
149         pvalues.append(result[1])
150
151
152     print "Dummy vs Perceptron p = ", pvalues[0]
153     print "Dummy vs Logistic p = ", pvalues[1]
154     print "Perceptron vs Logistic p = ", pvalues[2]
155
156
157     #part d: Standardization
158     scaler = preprocessing.StandardScaler().fit(train_data.X)
159     x_train_scaled = scaler.transform(train_data.X)
160     train_data_scaled = Data(X = x_train_scaled, y= train_data.y)
161
162     descriptive_stats_scaled = {'mean':[], 'stdev':[]}
163     score_array_scaled = []
164     for clf in clfs:
165         scores = cv_performance(clf, train_data_scaled, kfs)
166         mean_result = np.mean(scores)
167         stdev_result = np.std(scores)
168         descriptive_stats_scaled['mean'].append(mean_result)
169         descriptive_stats_scaled['stdev'].append(stdev_result)
170         score_array_scaled.append(scores)
171
172     pvalues_scaled = []
173     for combo in compare_t_tests:
174         result =
stats.ttest_rel(score_array_scaled[combo[0]].flatten(),
score_array_scaled[combo[1]].flatten())
175         pvalues_scaled.append(result[1])
176
177     print "Scaled Dummy vs Perceptron p = ", pvalues_scaled[0]
178     print "Scaled Dummy vs Logistic p = ", pvalues_scaled[1]
179     print "Scaled Perceptron vs Logistic p = ", pvalues_scaled[2]
180
181     #t-tests for the standardized vs. non-standardized
```

```
182
183     #Dummy Variable Row
184     pval_d = []
185     for i in range(3):
186         result = stats.ttest_rel(score_array_scaled[0].flatten(),
score_array[i].flatten())
187         pval_d.append(result[1])
188
189     print "Dummy Row vs Elements in Standardized"
190     print "D vs Standard D", pval_d[0]
191     print "D vs Standard P", pval_d[1]
192     print "D vs Standard L", pval_d[2]
193
194
195     #Perceptron Variable Row
196     pval_p = []
197     for i in range(3):
198         result = stats.ttest_rel(score_array_scaled[1].flatten(),
score_array[i].flatten())
199         pval_p.append(result[1])
200
201     print "Perceptron Row vs Elements in Standardized"
202     print "P vs Standard D", pval_p[0]
203     print "P vs Standard P", pval_p[1]
204     print "P vs Standard L", pval_p[2]
205
206     #Logistic Regression Variable Row
207     pval_r = []
208     for i in range(3):
209         result = stats.ttest_rel(score_array_scaled[2].flatten(),
score_array[i].flatten())
210         pval_r.append(result[1])
211
212     print "Logistic Row vs Elements in Standardized"
213     print "L vs Standard D", pval_r[0]
214     print "L vs Standard P", pval_r[1]
215     print "L vs Standard L", pval_r[2]
216
217     # print out descriptive_stats
218     print "Descriptive scaled stats in the order: Dummy, Perceptron,
Logistic"
219     print descriptive_stats_scaled
220
221
222     ### ===== TODO : END ===== ###
223
224     ### ===== TODO : START ===== ###
225     # part e: plot
226
227     # Indices of descriptive stats: 0 - Dummy, 1 - Perceptron, 2 -
```

Logistic

```

228
229     N = 2
230     ind = np.arange(N) # the x locations for the groups
231     width = 0.35       # the width of the bars
232     fig, ax = plt.subplots()
233
234
235     dum_mean = descriptive_stats['mean'][0]
236     dum_stdev = descriptive_stats['stdev'][0]
237
238     percep_means = (descriptive_stats['mean'][1],
239                     descriptive_stats_scaled['mean'][1])
240     percep_stdev = (descriptive_stats['stdev'][1],
241                    descriptive_stats_scaled['stdev'][1])
242     rects1 = ax.bar(ind + width, percep_means, width, color='r',
243 yerr=percep_stdev)
244
245     log_reg_means = (descriptive_stats['mean'][2],
246                     descriptive_stats_scaled['mean'][2])
247     log_reg_stdev = (descriptive_stats['stdev'][2],
248                     descriptive_stats_scaled['stdev'][2])
249     rects2 = ax.bar(ind, log_reg_means, width, color='b',
250 yerr=log_reg_stdev)
251
252     # add some text for labels, title and axes ticks
253     ax.set_ylabel('Accuracy')
254     ax.set_title('Accuracies by Preprocessing and Classifier')
255     ax.set_xticks(ind + width / 2)
256     ax.set_xticklabels(('No Preprocessing', 'Standardization'))
257
258     ax.legend((rects1[0], rects2[0]), ('Perceptron', 'Logistic
259 Regression'), loc='lower right')
260
261     def autolabel(rects):
262         """
263         Attach a text label above each bar displaying its height
264         """
265         for rect in rects:
266             height = rect.get_height()
267             ax.text(rect.get_x() + rect.get_width()/2., 1.05*height,
268                     '%.2f' % height,
269                     ha='center', va='bottom')
270
271
272     autolabel(rects1)
273     autolabel(rects2)

```

```
274
275     plt.axhline(y=dum_mean, color='k', linestyle='solid')
276     plt.axhline(y=(dum_mean + dum_stdev), color='k',
277                 linestyle='dashed')
278     plt.axhline(y=(dum_mean - dum_stdev), color='k',
279                 linestyle='dashed')
280
281     plt.show()
282
283     ### ===== TODO : END ===== ###
284
285 if __name__ == "__main__":
286     main()
```